



云计算上的图数据库 – Amazon Neptune

费良宏, Technical Evangelist @ AWS
lianghon@amazon.com

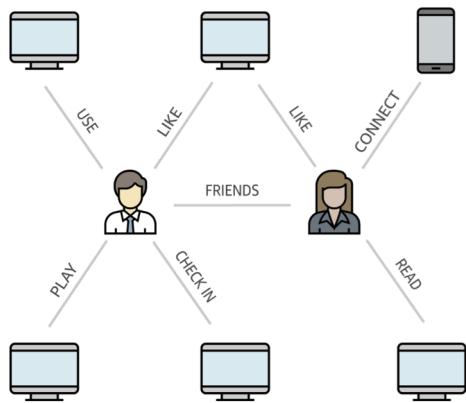


内容提要

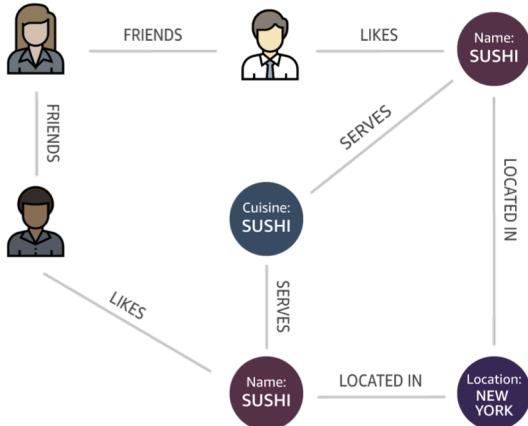
- 图数据库的特点及应用场景
- Amazon Neptune 概述
- Amazon Neptune 技术浅析
- 用户案例



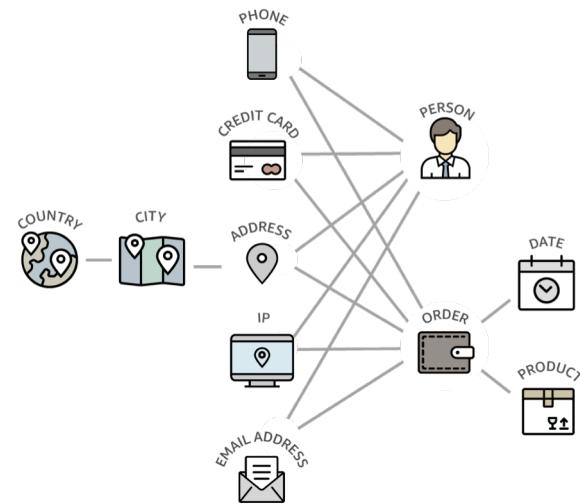
高度连接的数据



社交媒体



餐厅的推荐系统



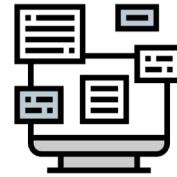
零售的欺诈检测



高度连接的数据的使用场景



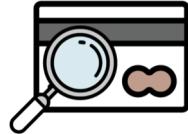
社交网络



推荐系统



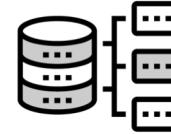
知识图谱



欺诈检测



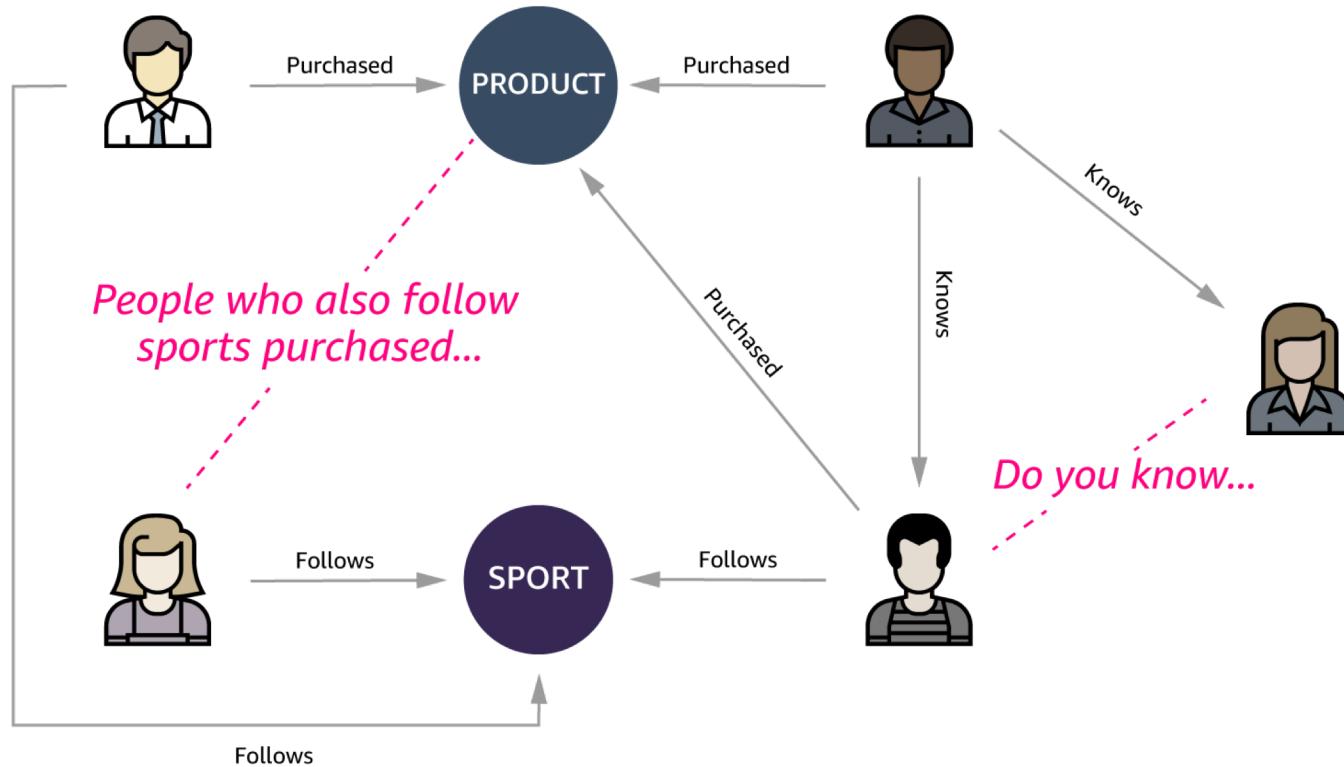
生命科学



网络 & IT 运维



基于关系的推荐系统

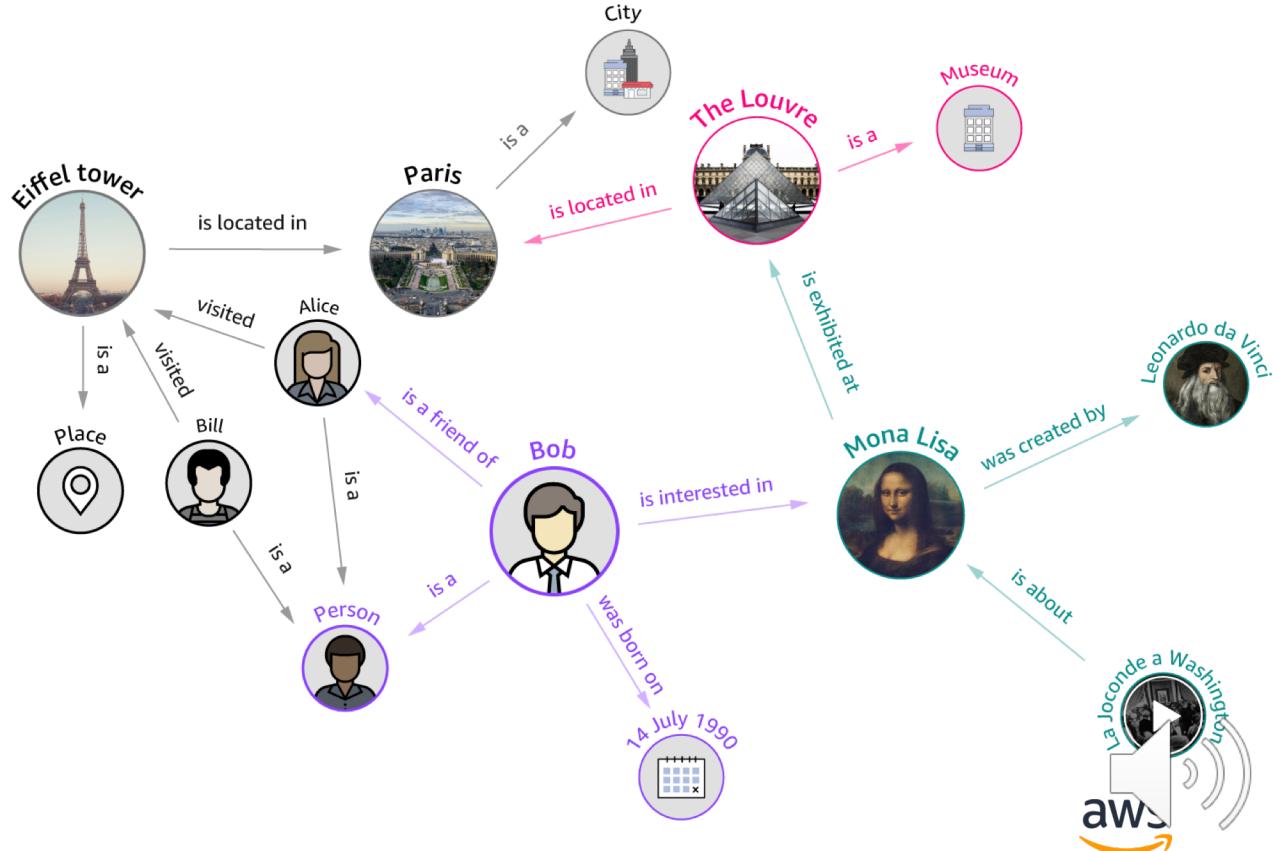


知识图谱应用

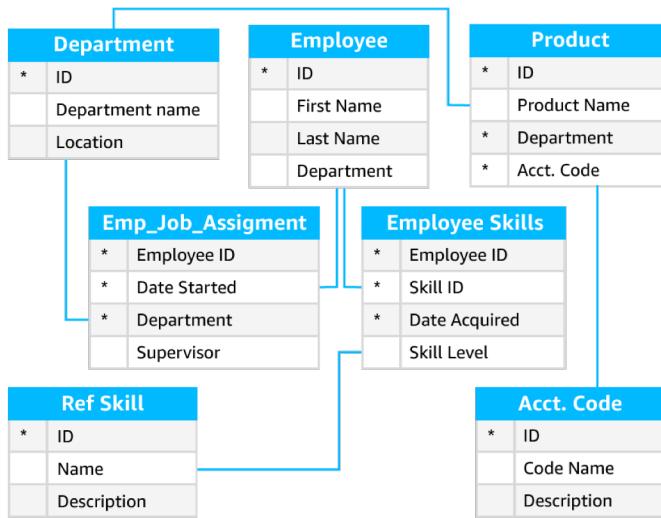
谁绘制了蒙娜丽莎？

爱丽丝在巴黎时应该参观哪些博物馆？

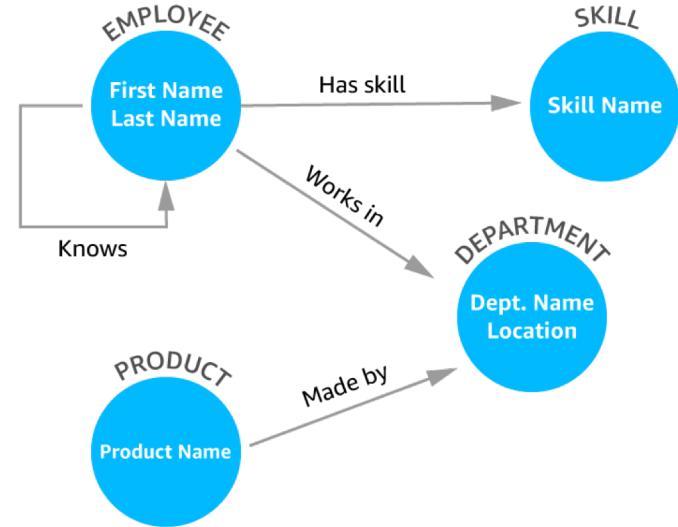
在卢浮宫有哪些艺术家的油画？



高度连接数据的处理方法



为特定业务流程而构建

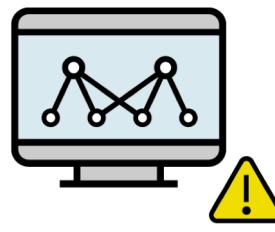


为回答关于关系的问题而构建

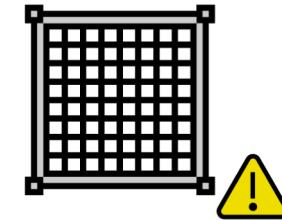
关系数据库的挑战 - 构建高度连接的数据应用



不自然的对图数据
的查询



低效的图处理



刚性模式对于数据的
变化是不灵活的

现有的图数据库的挑战



难于扩展



难以维护高可用性

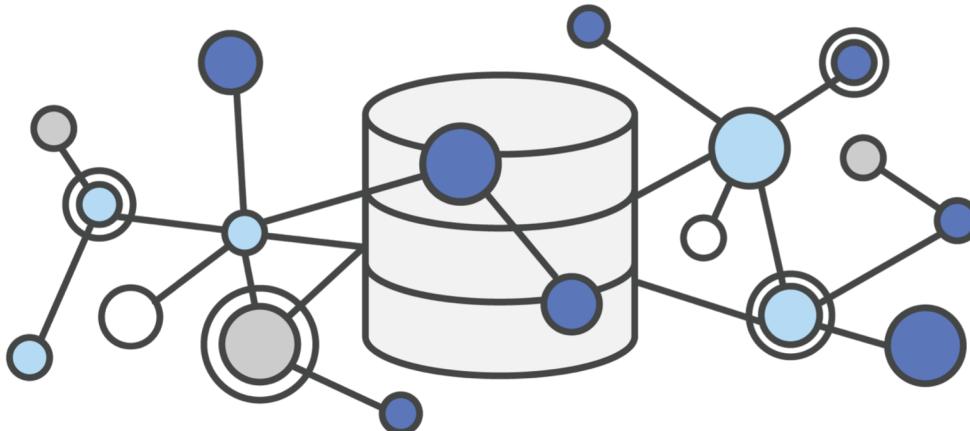


太昂贵



对开放标准的支持有限

图数据库 – 为高效存储和检索高度连接的数据而优化



AWS 上的数据库服务

Relational Databases



Amazon RDS



Amazon Redshift

Aurora

Commercial

Community

Data Warehouse



ORACLE[®]



PostgreSQL



PostgreSQL



Non-Relational Databases NEW!



Amazon
DynamoDB

Key Value

Document



Amazon
ElastiCache

In-Memory
Data Store



Amazon
Neptune

Graph



AWS Database Migration Service



图 (Graph) 模型和框架

属性图

(PROPERTY GRAPH)

开源的 Apache TinkerPop™

Gremlin Traversal Language



资源描述框架 RDF

(RESOURCE DESCRIPTION FRAMEWORK)

W3C 标准

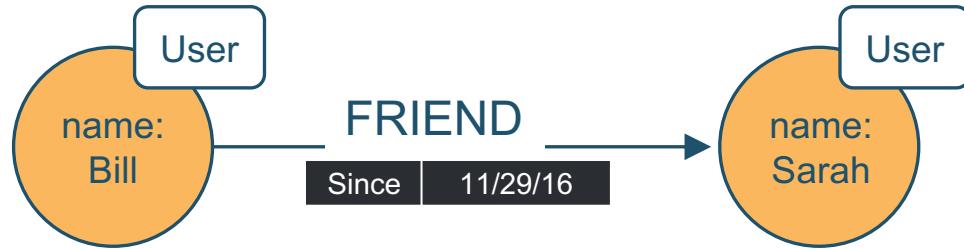
SPARQL Query Language



属性图

属性图是具有各自属性(即键/值对)的顶点和边的集合

- **顶点 (Vertex)** 顶点代表实体/域
- **边 (Edge)** 表示顶点之间的方向关系
 - 每条边都有一个表示关系类型的标签
- 每个顶点和边都有唯一的标识符
- 顶点和边可以有**属性**
- **属性**表示关于顶点和边的非关系信息



属性图 & APACHE TINKERPOP

- Apache TinkerPop

针对属性图的开源图形计算框架

- Gremlin

图遍历语言用于图的分析



Amazon Neptune 与 Tinkerpop Gremlin 3.3.0(最新版本于2017年8月发布)完全兼容，
并为Gremlin查询语言提供了优化的查询执行引擎

建立一个 TinkerPop 图

```
//Connect to Neptune and receive a remote graph, g.
```

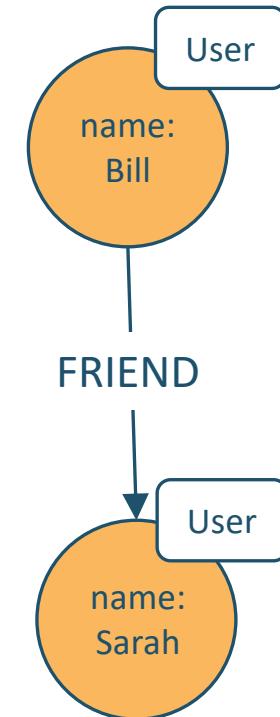
```
user1 = g.addVertex (id, 1, label, "User", "name", "Bill");  
user2 = g.addVertex (id, 2, label, "User", "name", "Sarah");
```

```
...
```

```
user1.addEdge("FRIEND", user2, id, 21);
```



Gremlin (Apache TinkerPop 3.3)



资源描述框架（RDF）图

- RDF图被描述为三元组的集合：主语、谓语和宾语

主语：→ <<http://www.socialnetwork.com/person#1>>
谓语：→ rdf:type contacts:User;
宾语：→ contact:name: "Bill" .



- 国际化资源标识符(IRIs)唯一地标识主题

IRI → <<http://www.socialnetwork.com/person#1>>

- 宾语可以是 IRI 或者文字

- RDF中的文字类似于属性，RDF支持XML数据类型
- 当对象是IRI时，它在图中形成一条“边”

主语：→ <<http://www.socialnetwork.com/person#1>>
谓语：→ contacts:friend
宾语 (IRI)：→ <<http://www.socialnetwork.com/person#2>> .



RDF 三元组示例

```
@prefix contacts: <http://www.socialnetwork.com/people#>.  
<http://www.socialnetwork.com/person#1>
```

```
    rdf:type contacts:User;  
    contact:name: "Bill" .
```

```
<http://www.socialnetwork.com/person#1>  
    contacts:friend <http://www.socialnetwork.com/person#2> .
```

```
<http://www.socialnetwork.com/person#2>  
    rdf:type contacts:User;  
    contact:name: "Sarah" .
```



RDF
(Turtle Serialization)

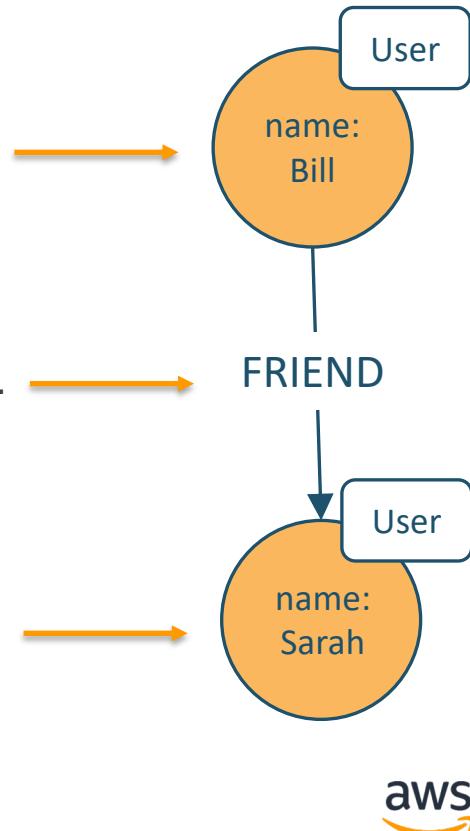
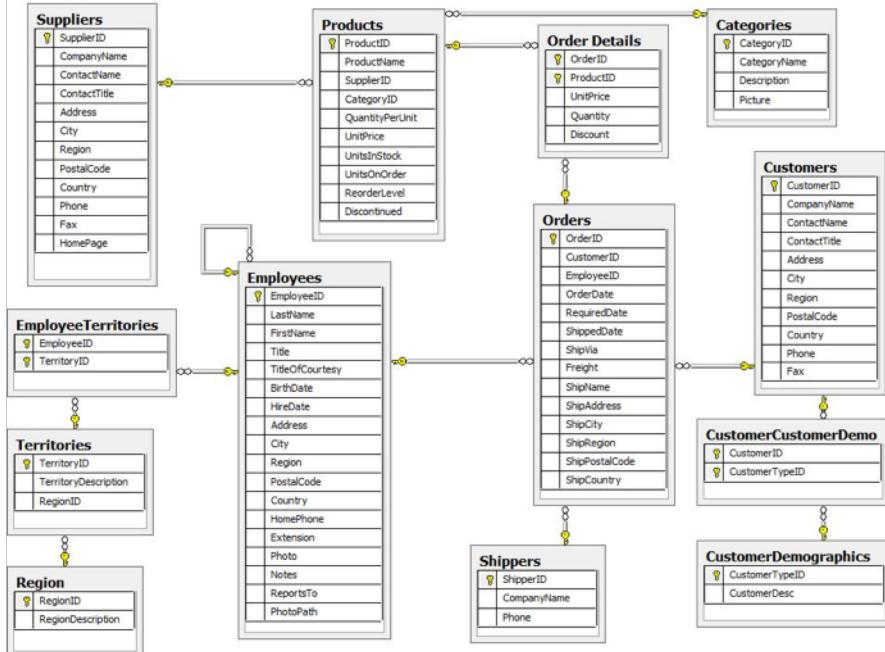
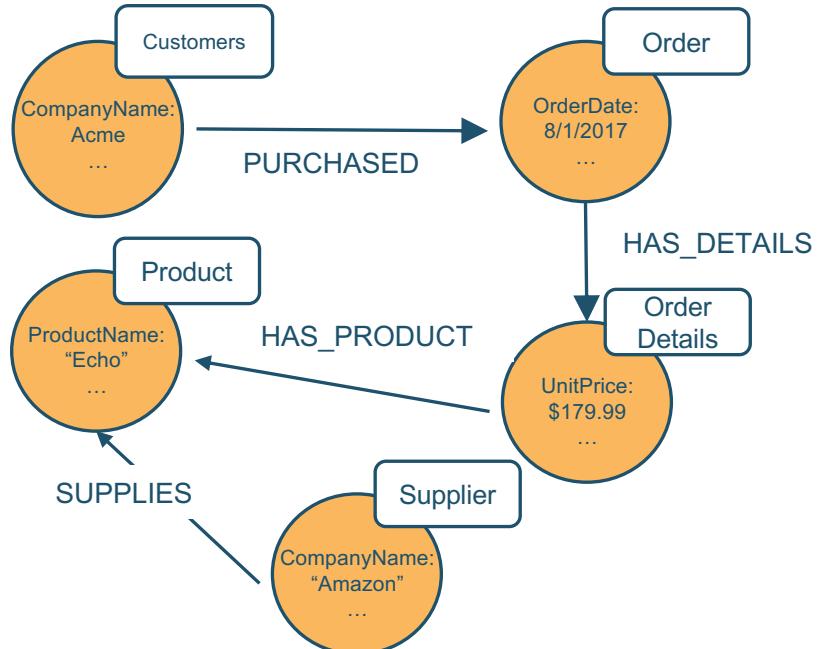


图 vs. 关系数据库模型

关系模型



图模型子集



* 来源：<http://www.playnextrac.com/index.html#show:article>



SQL 关系数据库查询

查找购买过Echo产品的公司的名称

```
SELECT distinct c.CompanyName
FROM customers AS c
    JOIN orders AS o          /* Join the customer from the order */
        (c.CustomerID = o.CustomerID)
    JOIN order_details AS od   /* Join the order details from the order */
/*
        ON (o.OrderID = od.OrderID)
    JOIN products as p         /* Join the products from the order details
/*
        ON (od.ProductID = p.ProductID)
WHERE p.ProductName = 'Echo'; /* Find the product named 'Echo' */
```



SPARQL 声明式图查询

查找购买过Echo产品的公司的名称

```
PREFIX sales_db: <http://sales.widget.com/>
SELECT distinct ?comp_name WHERE {
    ?customer <sales_db:HAS_ORDER> ?order ;           #customer graph pattern
              <sales_db:CompanyName> ?comp_name .         #orders graph pattern
    ?order   <sales_db:HAS_DETAILS> ?order_d .        #order details graph pattern
    ?order_d <sales_db:HAS_PRODUCT> ?product .       #products graph
pattern
    ?product  <sales_db:ProductName> "Echo" .
}
```

* 来源：<http://www.playnexus.com/index.html#show:article>



Gremlin 命令式图遍历

查找购买过Echo产品的公司的名称

```
/* All products named "Echo" */
g.v().hasLabel('Product').has('name', 'Echo')
    .in('HAS_PRODUCT') /* Traverse to order details */
    .in('HAS_DETAILS') /* Traverse to order */
    .in('HAS_ORDER') /* Traverse to Customer */
    .values('CompanyName').dedup() /* Unique Company Name */
```



Python 样例代码（Gremlin）

- 安装 gremlinpython 程序包

```
pip3 install gremlinpython --user
```

- 遍历图

```
from gremlin_python import statics
from gremlin_python.structure.graph import Graph
from gremlin_python.process.graph_traversal import __
from gremlin_python.process.strategies import *
from gremlin_python.driver.driver_remote_connection import DriverRemoteConnection

graph = Graph()
g = graph.traversal().withRemote(DriverRemoteConnection('ws://your-neptune-endpoint:8182/gremlin','g'))
print(g.V().limit(2).toList())
```



AMAZON NEPTUNE

全托管的图数据库

开放



支持 Apache TinkerPop
& W3C RDF 图模型

快速



查询十亿级关系数据
仅有毫秒级延迟

可靠



数据拥有跨3个可用区
的6个副本，具有完全
备份和恢复

易用



使用Gremlin和SPARQL
轻松构建强大的查询



全托管的服务

优点

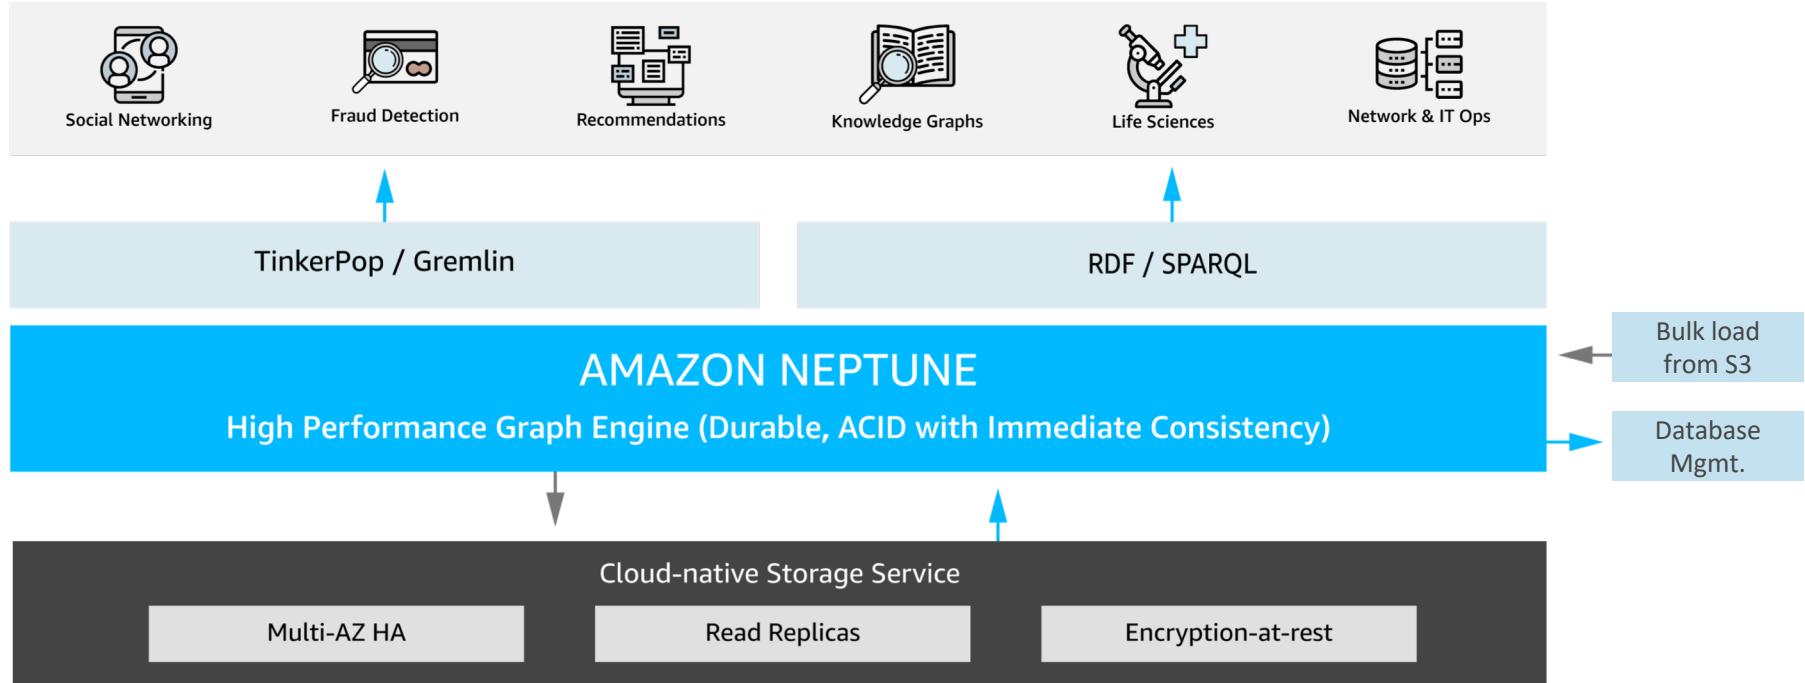
- ✓ 通过控制台容易配置
- ✓ Multi-AZ高可用性、ACID
- ✓ 支持最多15个只读副本
- ✓ 支持数据加密存储
- ✓ 支持传输加密(TLS)
- ✓ 备份和恢复，时间点恢复

The screenshot shows the AWS Lambda service page. At the top, there's a search bar and navigation links for 'Services' and 'Resource Groups'. Below the search bar, the Lambda logo is displayed with the text 'Lambda Functions'. On the left, a sidebar menu includes 'Functions', 'Triggers', 'Logs', 'Metrics', and 'CloudWatch Metrics'. The main content area shows a table with one row for the function 'hello'. The table columns are 'Function Name', 'Last Run', 'Status', and 'Actions'. The 'hello' function has a status of 'Running' and an 'Edit' button.

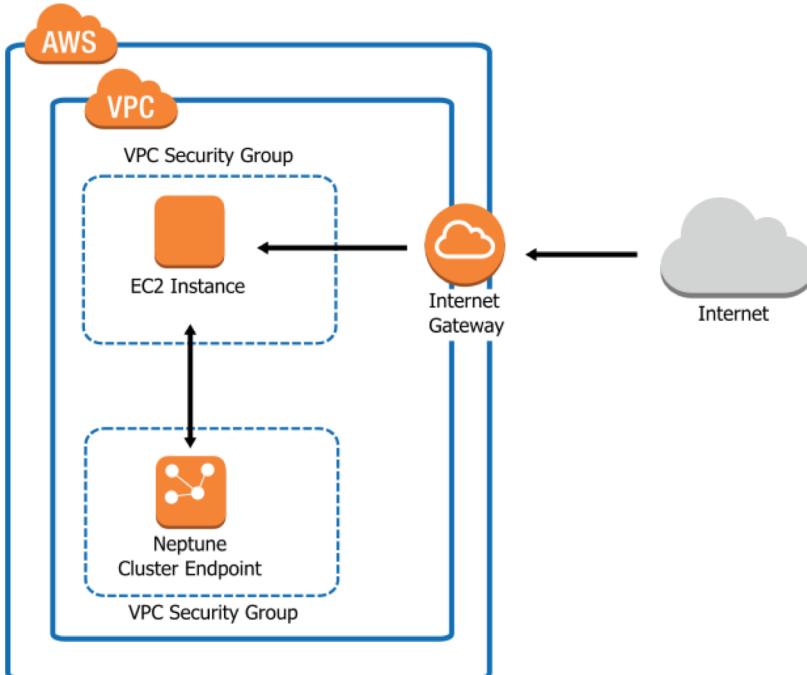
The screenshot shows the AWS Neptune service page. At the top, there's a search bar and navigation links for 'Services' and 'Resource Groups'. Below the search bar, the Neptune logo is displayed with the text 'Graph Database'. The main content area shows a section titled 'How it works' with a brief description of Neptune's capabilities. On the left, a sidebar menu includes 'Clusters', 'Instances' (which is highlighted in orange), 'Snapshots', 'Security groups', 'Subnet groups', 'Parameter groups', 'Events', and 'Event subscriptions'. The main content area shows a table titled 'Instances (11)' with a search bar above it. The table columns are 'DB Instance', 'Engine', and 'Status'. There are three rows listed: 'neptune-beta-01', 'neptune-beta-02', and 'neptune-beta-03', all of which are in the 'available' status.



AMAZON NEPTUNE 架构



AMAZON NEPTUNE: VPC 部署



- 安全部署在VPC
- 通过在两个不同可用区(AZs)的两个子网中部署来增加可用性
- 集群组总是跨越三个AZ以提供持久性的存储
- 有关VPC安装细节, 请参阅[Amazon Neptune 文档](#)

Amazon Neptune 存储引擎概述

在3个可用区中数据被复制6份

数据持续备份到Amazon S3
(耐久性为11个9)

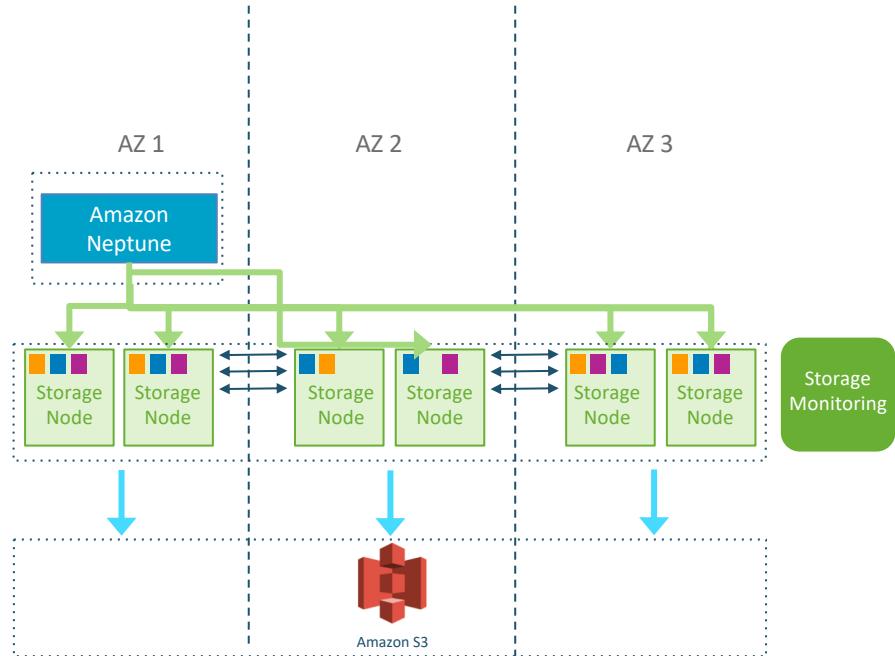
持续监控节点和磁盘的修复

10 GB段作为修复单位或平衡的热点

法定读/写系统的数量; 容忍延迟

仲裁成员资格的变化不会延迟写入

存储容量自动增长到64 TB



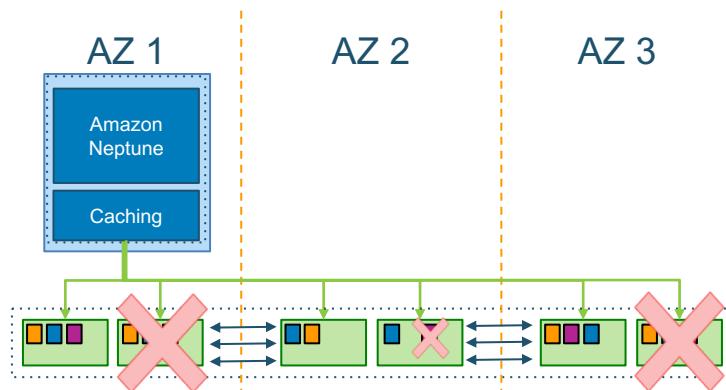
Amazon Neptune 高可用性和容错能力

为何失效?

段故障(磁盘)

节点故障(机器)

可用区故障(网络或数据中心)

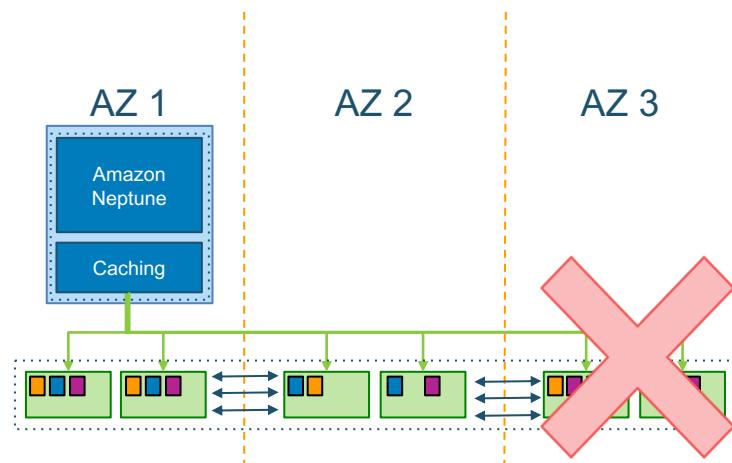


优化方法:

6个副本中4个写成功是必要的

6个副本中4个读成功是必要的

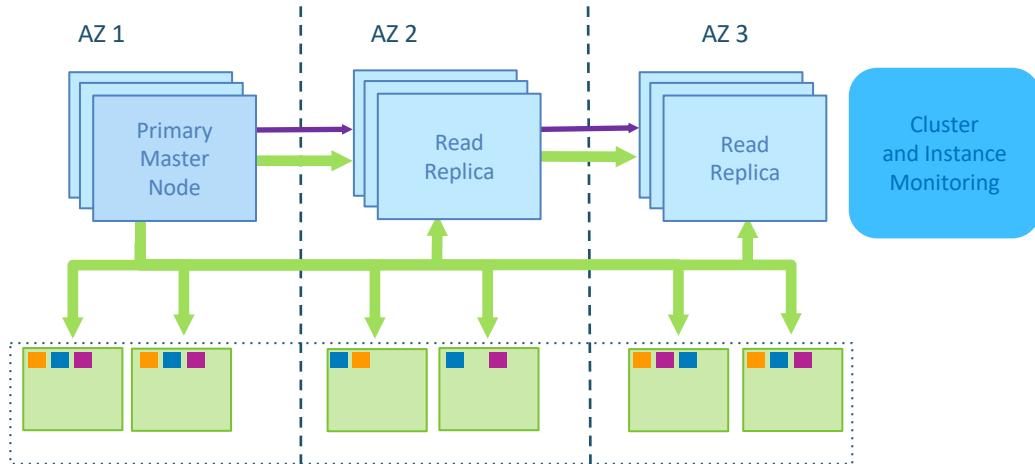
对等复制以修复故障



Amazon Neptune 读复制

可用性

- 失败的数据库节点将被自动检测和替换
- 失败的数据库进程将被自动检测和回收
- 如果需要，副本将自动提升为主副本(故障转移)
- 客户指定的故障转移

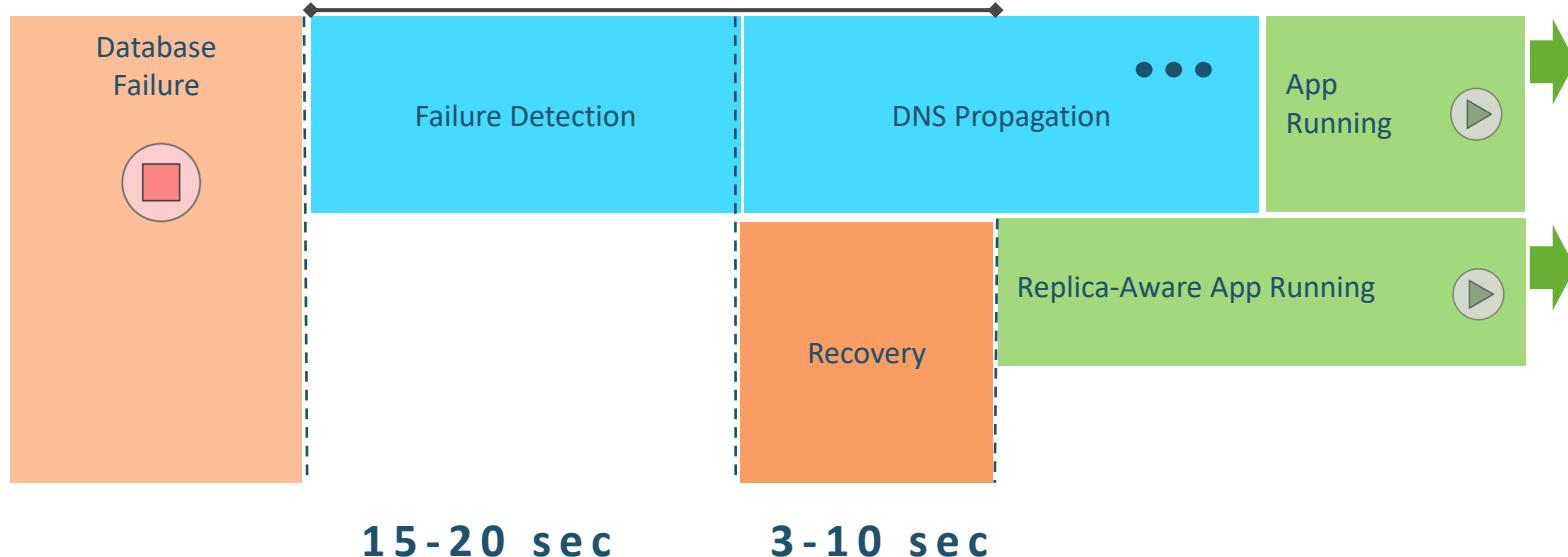


性能

- 客户应用程序可以跨读副本扩展读流量
- 在读副本之间进行读操作的平衡



Amazon Neptune故障转移时间通常小于30秒



AMAZON NEPTUNE 的数据加载

Gremlin 加载数据格式

- 必须采用UTF-8格式编码
- CSV 格式遵循 RFC 4180 CSV 规范



RDF 加载数据格式

- 必须采用UTF-8格式编码
- 规范 (<https://www.w3.org/TR/n-triples/>) 中的 N-Triples (ntriples)
- 规范 (<https://www.w3.org/TR/n-quads/>) 中的 N-Quads (nquads)
- 规范 (<https://www.w3.org/TR/rdf-syntax-grammar/>) 中的 RDF/XML (rdfxml)
- 规范 (<https://www.w3.org/TR/turtle/>) 中的 N-Quads (turtle)

示例：加载数据

```
curl -X POST \
-H 'Content-Type: application/json' \
http://your-neptune-endpoint:8182/loader -d '
{
  "source" : "s3://bucket-name/object-key-name",
  "format" : "format",
  "iamRoleArn" : "arn:aws:iam::account-id:role/role-name",
  "region" : "region",
  "failOnError" : "FALSE"
}'
```



浏览全球税收政策的网络



THOMSON REUTERS

“我们的客户越来越需要浏览复杂的全球税收政策和法规网络。我们需要一种方法来模拟我们最大客户的复杂企业结构，并提供端到端税收解决方案。我们的平台使用微服务体系结构，并开始利用Amazon Neptune作为基于图的系统，快速在数据中创建链接。”

- Tim Vanderham, CTO , Thomson Reuters Tax & Accounting



用户案例



SIEMENS



AstraZeneca The AstraZeneca logo, featuring the company name in purple and a small orange and yellow molecular or leaf-like icon to the right.



intuit.



SAMSUNG

Blackfynn

The Blackfynn logo, featuring a series of blue horizontal bars of varying lengths followed by the word "Blackfynn" in a sans-serif font.

AMAZON NEPTUNE

全托管的图数据库

开放



支持 Apache TinkerPop
& W3C RDF 图模型

快速



查询十亿级关系数据
仅有毫秒级延迟

可靠



数据拥有跨3个可用区
的6个副本，具有完全
备份和恢复

易用



使用Gremlin和SPARQL
轻松构建强大的查询



AMAZON NEPTUNE 的资源

- [Amazon Neptune 服务介绍](#)
- [Amazon Neptune 文档](#)
- [Amazon Neptune Tools Repo](#)
- [Amazon Neptune Samples Repo](#)

README.md

GraphML 2 Neptune CSV

This Python script provides a utility to convert GraphML files into the CSV format that is used by Amazon Neptune for Bulk Loading.

Usage

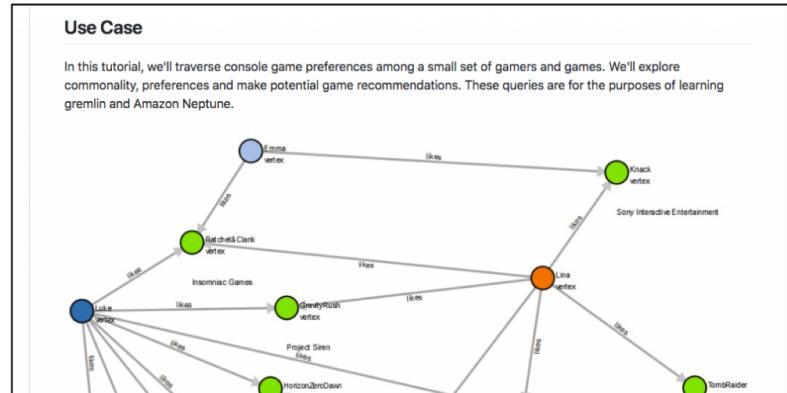
```
Usage: graphml2csv.py [options]
```

Copyright 2018 Amazon.com, Inc. or its affiliates.
Licensed under the Apache License 2.0
<http://aws.amazon.com/apache2.0/>

Options:

```
--version           show program's version number and exit
--help              show this help message and exit
--i FILE            set input path [default: none]
--v, --verbose      set verbosity level [default: none]
```

A utility python script to convert GraphML files into the Amazon Neptune CSV format for bulk ingestion. See <https://docs.aws.amazon.com/neptune/latest/userguide/bulk-load-tutorial-format-gremlin.html>.



感谢！